

Import packages

```
In [1]: import pandas as pd
import numpy as np
import scipy.stats as stats
import matplotlib.pyplot as plt
import seaborn as sns
import plotly.express as px
```

```
In [2]: # check data directory
# ls
```

import data into pandas dataframe

```
In [3]: medical_df = pd.read_csv("project_dataset - project_dataset.csv")
medical_df.head(10)
```

```
Out[3]:
```

	EmployeeID	Age	BusinessTravel	MonthlyIncome	JobSatisfaction	Bonus	Department	DistanceFromHome	Education	EducationField	EnvSatisfaction	Gender	JobRole	MaritalStatus	PerformanceRating	TrainingTir
0	1	41	Travel_Rarely	5993	4	17979	Sales	1	2	Life Sciences	2	Female	Sales Executive	Single	3	
1	2	49	Travel_Frequently	5130	2	20520	Research & Development	8	1	Life Sciences	3	Male	Research Scientist	Married	4	
2	4	37	Travel_Rarely	2090	3	6270	Research & Development	2	2	Other	4	Male	Laboratory Technician	Single	3	
3	5	33	Travel_Frequently	2909	3	8727	Research & Development	3	4	Life Sciences	4	Female	Research Scientist	Married	3	
4	7	27	Travel_Rarely	3468	2	10404	Research & Development	2	1	Medical	1	Male	Laboratory Technician	Married	3	
5	8	32	Travel_Frequently	3068	4	9204	Research & Development	2	2	Life Sciences	4	Male	Laboratory Technician	Single	3	
6	10	59	Travel_Rarely	2670	1	10680	Research & Development	3	3	Medical	3	Female	Laboratory Technician	Married	4	
7	11	30	Travel_Rarely	2693	3	10772	Research & Development	24	1	Life Sciences	4	Male	Laboratory Technician	Divorced	4	
8	12	38	Travel_Frequently	9526	3	38104	Research & Development	23	3	Life Sciences	4	Male	Manufacturing Director	Single	4	
9	13	36	Travel_Rarely	5237	3	15711	Research & Development	27	3	Medical	3	Male	Healthcare Representative	Married	3	

```
In [4]: # make the EmployeeID as the index column
medical_df = medical_df.set_index('EmployeeID')
medical_df
```

Out[4]:

	Age	BusinessTravel	MonthlyIncome	JobSatisfaction	Bonus	Department	DistanceFromHome	Education	EducationField	EnvSatisfaction	Gender	JobRole	MaritalStatus	PerformanceRating	TrainingTimes1
EmployeeID															
1	41	Travel_Rarely	5993	4	17979	Sales	1	2	Life Sciences	2	Female	Sales Executive	Single	3	
2	49	Travel_Frequently	5130	2	20520	Research & Development	8	1	Life Sciences	3	Male	Research Scientist	Married	4	
4	37	Travel_Rarely	2090	3	6270	Research & Development	2	2	Other	4	Male	Laboratory Technician	Single	3	
5	33	Travel_Frequently	2909	3	8727	Research & Development	3	4	Life Sciences	4	Female	Research Scientist	Married	3	
7	27	Travel_Rarely	3468	2	10404	Research & Development	2	1	Medical	1	Male	Laboratory Technician	Married	3	
...
2061	36	Travel_Frequently	2571	4	7713	Research & Development	23	2	Medical	3	Male	Laboratory Technician	Married	3	
2062	39	Travel_Rarely	9991	1	29973	Research & Development	6	1	Medical	4	Male	Healthcare Representative	Married	3	
2064	27	Travel_Rarely	6142	2	24568	Research & Development	4	3	Life Sciences	2	Male	Manufacturing Director	Married	4	
2065	49	Travel_Frequently	5390	2	16170	Sales	2	3	Medical	4	Male	Sales Executive	Married	3	
2068	34	Travel_Rarely	4404	3	13212	Research & Development	8	3	Medical	2	Male	Laboratory Technician	Married	3	

1470 rows × 19 columns

Data Preparation

Perform Exploration Data Analysis

- check the data shape
- Check for Null Data
- Check for Duplicates
- Check the data summary statistics
- Check for correlations
- Check for multicollinearity
- Rearranging the columns

In [5]: `# check the data shape`
`medical_df.shape`

Out[5]: (1470, 19)

In [6]: `# check for null data`
`medical_df.isnull().sum()`

```
Out[6]: Age                0
        BusinessTravel    0
        MonthlyIncome     0
        JobSatisfaction    0
        Bonus             0
        Department        0
        DistanceFromHome  0
        Education         0
        EducationField     0
        EnvSatisfaction    0
        Gender            0
        JobRole           0
        MaritalStatus     0
        PerformanceRating  0
        TrainingTimesLastYear 0
        YearsAtCompany    0
        YearsSinceLastPromotion 0
        OverTime          0
        Attrition         0
        dtype: int64
```

```
In [7]: #check for duplicates
        medical_df[medical_df.duplicated(keep=False)]
```


```
Out[7]:      Age  BusinessTravel  MonthlyIncome  JobSatisfaction  Bonus  Department  DistanceFromHome  Education  EducationField  EnvSatisfaction  Gender  JobRole  MaritalStatus  PerformanceRating  TrainingTimesLastYear
EmployeeID
```

```
In [8]: #check the data summary statistics
        medical_df.describe(include='all')
```

```
Out[8]:
```

	Age	BusinessTravel	MonthlyIncome	JobSatisfaction	Bonus	Department	DistanceFromHome	Education	EducationField	EnvSatisfaction	Gender	JobRole	MaritalStatus	PerformanceRating	TrainingTi
count	1470.000000	1470	1470.000000	1470.000000	1470.000000	1470	1470.000000	1470.000000	1470	1470.000000	1470	1470	1470	1470.000000	
unique	NaN	3	NaN	NaN	NaN	3	NaN	NaN	6	NaN	2	9	3	NaN	
top	NaN	Travel_Rarely	NaN	NaN	NaN	Research & Development	NaN	NaN	Life Sciences	NaN	Male	Sales Executive	Married	NaN	
freq	NaN	1043	NaN	NaN	NaN	961	NaN	NaN	606	NaN	882	326	673	NaN	
mean	36.923810	NaN	6502.931293	2.728571	20479.501361	NaN	9.192517	2.912925	NaN	2.721769	NaN	NaN	NaN	3.153741	
std	9.135373	NaN	4707.956783	1.102846	15066.272964	NaN	8.106864	1.024165	NaN	1.093082	NaN	NaN	NaN	0.360824	
min	18.000000	NaN	1009.000000	1.000000	3027.000000	NaN	1.000000	1.000000	NaN	1.000000	NaN	NaN	NaN	3.000000	
25%	30.000000	NaN	2911.000000	2.000000	9333.750000	NaN	2.000000	2.000000	NaN	2.000000	NaN	NaN	NaN	3.000000	
50%	36.000000	NaN	4919.000000	3.000000	15484.500000	NaN	7.000000	3.000000	NaN	3.000000	NaN	NaN	NaN	3.000000	
75%	43.000000	NaN	8379.000000	4.000000	26103.750000	NaN	14.000000	4.000000	NaN	4.000000	NaN	NaN	NaN	3.000000	
max	60.000000	NaN	19999.000000	4.000000	79892.000000	NaN	29.000000	5.000000	NaN	4.000000	NaN	NaN	NaN	4.000000	

```
In [9]: #use heatmap function to understand the data summary better
        # use seaborn to visualize the correlation heatmap
        plt.figure(figsize=(8,8))
        sns.heatmap(medical_df.corr().round(2),vmin=-1, vmax=1, annot=True,linewidth=.5);
```



```
In [10]: medical_df.columns
```

```
Out[10]: Index(['Age', 'BusinessTravel', 'MonthlyIncome', 'JobSatisfaction', 'Bonus',
      'Department', 'DistanceFromHome', 'Education', 'EducationField',
      'EnvSatisfaction', 'Gender', 'JobRole', 'MaritalStatus',
      'PerformanceRating', 'TrainingTimesLastYear', 'YearsAtCompany',
      'YearsSinceLastPromotion', 'OverTime', 'Attrition'],
      dtype='object')
```

```
In [11]: # check for multicollinearity
from statsmodels.stats.outliers_influence import variance_inflation_factor
copy_medical = medical_df.copy()
copy_medical['Attrition'] = copy_medical['Attrition'].map({'Yes':1,'No':0})
copy_medical['OverTime'] = copy_medical['OverTime'].map({'Yes':1,'No':0})
copy_medical['MaritalStatus'] = copy_medical['MaritalStatus'].map({'Married':1,'Divorced':2,'Single':3})
copy_medical['BusinessTravel'] = copy_medical['BusinessTravel'].map({'Travel_Rarely':1,'Travel_Frequently':2,'Non-Travel':3})
copy_medical['Gender'] = copy_medical['Gender'].map({'Female':1,'Male':0})

# select certain input variables
copy_medical = copy_medical[['Age', 'BusinessTravel', 'MonthlyIncome', 'JobSatisfaction', 'Bonus',
      'DistanceFromHome',
      'EnvSatisfaction', 'Gender', 'MaritalStatus',
      'PerformanceRating', 'TrainingTimesLastYear', 'YearsAtCompany',
      'YearsSinceLastPromotion', 'OverTime', 'Attrition']]

copy_medical
```

Out[11]:

	Age	BusinessTravel	MonthlyIncome	JobSatisfaction	Bonus	DistanceFromHome	EnvSatisfaction	Gender	MaritalStatus	PerformanceRating	TrainingTimesLastYear	YearsAtCompany	YearsSinceLastPromotion	OverTime
EmployeeID														
1	41	1	5993	4	17979	1	2	1	3	3	0	6	0	
2	49	2	5130	2	20520	8	3	0	1	4	3	10	1	
4	37	1	2090	3	6270	2	4	0	3	3	3	0	0	
5	33	2	2909	3	8727	3	4	1	1	3	3	8	3	
7	27	1	3468	2	10404	2	1	0	1	3	3	2	2	
...	
2061	36	2	2571	4	7713	23	3	0	1	3	3	5	0	
2062	39	1	9991	1	29973	6	4	0	1	3	5	7	1	
2064	27	1	6142	2	24568	4	2	0	1	4	0	6	0	
2065	49	2	5390	2	16170	2	4	0	1	3	3	9	0	
2068	34	1	4404	3	13212	8	2	0	1	3	3	4	1	

1470 rows × 15 columns

In [12]:

```
# VIF Dataframe
vif_data = pd.DataFrame()
vif_data['feature'] = copy_medical.columns
vif_data
```

Out[12]:

	feature
0	Age
1	BusinessTravel
2	MonthlyIncome
3	JobSatisfaction
4	Bonus
5	DistanceFromHome
6	EnvSatisfaction
7	Gender
8	MaritalStatus
9	PerformanceRating
10	TrainingTimesLastYear
11	YearsAtCompany
12	YearsSinceLastPromotion
13	OverTime
14	Attrition

```
In [13]: #calculating VIF for each feature
vif_data['VIF'] = [variance_inflation_factor(copy_medical.values,i) for i in range(len(copy_medical.columns))]
print(vif_data)
```

	feature	VIF
0	Age	22.295740
1	BusinessTravel	5.312998
2	MonthlyIncome	111.293878
3	JobSatisfaction	7.053365
4	Bonus	106.448691
5	DistanceFromHome	2.302447
6	EnvSatisfaction	7.161742
7	Gender	1.678806
8	MaritalStatus	5.674453
9	PerformanceRating	40.294776
10	TrainingTimesLastYear	5.669428
11	YearsAtCompany	4.540620
12	YearsSinceLastPromotion	2.392605
13	OverTime	1.526281
14	Attrition	1.409794

```
In [14]: # copy_medical.columns
```

```
In [15]: # copy_medical = copy_medical[['Age', 'DistanceFromHome', 'MonthlyIncome', 'JobSatisfaction', 'EnvSatisfaction', 'YearsAtCompany', 'Gender', 'OverTime', 'Attrition']]
```

```
In [16]: # copy_medical = pd.get_dummies(copy_medical)
# copy_medical
```

```
In [17]: # plt.figure(figsize=(8,8))
# sns.heatmap(copy_medical.corr().round(2),vmin=-1, vmax=1, annot=True);
```


Check for variables that affect the attrition rate

Univariate Analysis

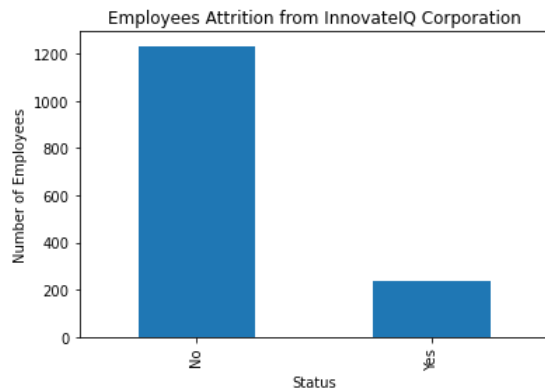
```
In [18]: # use value count function to find the number of counts of
# attributes within the class
# and use normalize to find the proportion of it
attrition = medical_df['Attrition'].value_counts(normalize=True)*100
attrition
```

```
Out[18]: No      83.877551
Yes      16.122449
Name: Attrition, dtype: float64
```

```
In [19]: attrition.plot(kind='pie', autopct='%0.1f%%')
plt.show()
```



```
In [20]: # use graphical to understand the data better
medical_raw_count = medical_df['Attrition'].value_counts()
medical_raw_count.plot(kind='bar')
plt.xlabel('Status')
plt.ylabel('Number of Employees')
plt.title('Employees Attrition from InnovateIQ Corporation')
plt.show()
```




Finding For Attrition:

The data for attrition column shows about 84% of their client are still working in the company.while about 16% has attrited from the medical company.However the class proportion is unbalanced.

```
In [21]: gender = medical_df['Gender'].value_counts(normalize=True)*100
gender
```

```
Out[21]: Male      60.0
Female    40.0
Name: Gender, dtype: float64
```

```
In [22]: gender.plot(kind='pie', autopct='%1f%%')
plt.show()
```



Finding For Gender:

The data for gender column shows that 60% of their employees are male and 40% are female regardless their attrition status


```
In [23]: years_at_com = medical_df['YearsAtCompany'].describe()
years_at_com
```

```
Out[23]: count    1470.000000
mean       7.008163
std        6.126525
min        0.000000
25%        3.000000
50%        5.000000
75%        9.000000
max        40.000000
Name: YearsAtCompany, dtype: float64
```


```
In [24]: medical_df['YearsAtCompany'].plot(kind='hist',bins=40)
plt.title('Working Years Distribution in Organization')
# check for the skewness and kurtosis
print("Skewness: {:.3f}".format(medical_df['YearsAtCompany'].skew()))
print("Kurtosis: {:.3f}".format(medical_df['YearsAtCompany'].kurt()))
```

Skewness: 1.765

Kurtosis: 3.936




```
In [25]: sns.boxplot(x='YearsAtCompany',data=medical_df)
plt.show()
```


In [26]: `medical_df['Age'].describe()`

```
Out[26]: count    1470.000000
mean      36.923810
std        9.135373
min       18.000000
25%       30.000000
50%       36.000000
75%       43.000000
max       60.000000
Name: Age, dtype: float64
```

In [27]: `medical_df['Age'].plot(kind='hist',bins=40)`
`plt.xlabel('Age')`
`plt.ylabel('Number of Employees')`
`plt.title('Age Distribution in Organization')`
`plt.show()`




In [28]: `sns.boxplot(x='Age',data=medical_df)`
`plt.show()`




```
In [29]: medical_df['MonthlyIncome'].plot(kind='hist',bins=20)
```

```
Out[29]: <AxesSubplot:ylabel='Frequency'>
```




```
In [30]: sns.boxplot(x='MonthlyIncome',data=medical_df)
plt.show()
```




```
In [31]: medical_df['PerformanceRating'].plot(kind='hist')
```

```
Out[31]: <AxesSubplot:ylabel='Frequency'>
```



```
In [32]: envi = medical_df['EnvSatisfaction'].plot(kind='hist')
envi
```

```
Out[32]: <AxesSubplot:ylabel='Frequency'>
```




```
In [33]: envi_prop = medical_df['EnvSatisfaction'].value_counts(normalize=True)*100
envi_prop
```

```
Out[33]: 3    30.816327
4    30.340136
2    19.523810
1    19.319728
Name: EnvSatisfaction, dtype: float64
```


```
In [34]: envi_prop.plot(kind='pie', autopct='%0.1f%%')
```

```
Out[34]: <AxesSubplot:ylabel='EnvSatisfaction'>
```




```
In [35]: promote = medical_df['YearsSinceLastPromotion'].plot(kind='hist')
promote
```

```
Out[35]: <AxesSubplot:ylabel='Frequency'>
```




```
In [36]: sns.boxplot(x='YearsSinceLastPromotion', data=medical_df)
plt.show()
```



```
In [37]: over_prop = medical_df['OverTime'].value_counts(normalize=True)*100
over_prop
```

```
Out[37]: No      71.70068
Yes       28.29932
Name: OverTime, dtype: float64
```

```
In [38]: over_prop.plot(kind='pie', autopct='%1f%%')
plt.title('Employees Worked Overtime')
plt.show()
```



Multivariate Analysis

Describe the Attrition based on Gender. Compute the attrition rate based on Gender

```
In [39]: round(pd.crosstab(index=medical_df['Attrition'], columns=medical_df['Gender'], normalize='index'), 2)
```

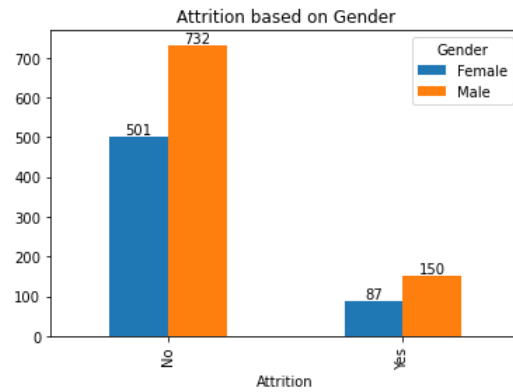
```
Out[39]:
```

Gender	Female	Male
No	0.41	0.59
Yes	0.37	0.63

Attrition		
No	0.41	0.59
Yes	0.37	0.63

```
In [40]: gender = pd.crosstab(index=medical_df['Attrition'], columns=medical_df['Gender'])
ax = gender.plot(kind = 'bar')

# annotate
ax.bar_label(ax.containers[0], label_type='edge')
ax.bar_label(ax.containers[1], label_type='edge')
plt.title('Attrition based on Gender')
plt.show()
```



```
In [41]: # test with CHI SQUARE test to check if proportion different for female and male by Attrition status has differently significant
chi,p,dof,exp = stats.chi2_contingency(gender)
print('the chi square test is {}'.format(round(p,3)))
# p > 0.05
#The ratios in our data are the same. Gender isn't a strong determinant of an employee being attrited or not
```

the chi square test is 0.291

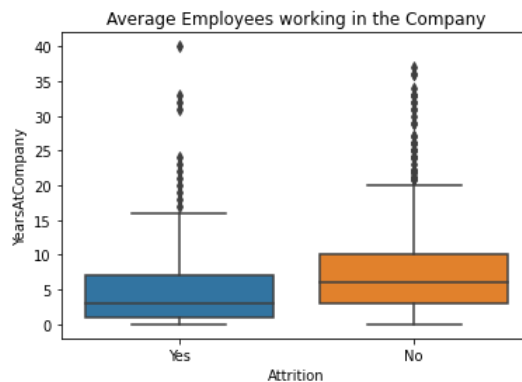
Describe the Attrition based on YearsAtCompany. Compute the attrition rate based on YearsAtCompany

```
In [42]: mean_years = round(medical_df.groupby('Attrition',as_index=False)['YearsAtCompany'].mean())
mean_years = mean_years.rename(columns={'YearsAtCompany': 'Mean Years working at the Company'})
mean_years
```


```
Out[42]:
```

	Attrition	Mean Years working at the Company
0	No	7.0
1	Yes	5.0

```
In [43]: sns.boxplot(x='Attrition', y='YearsAtCompany', data = medical_df)
plt.title('Average Employees working in the Company')
plt.show()
```



```
In [44]: fig = plt.figure(figsize=(18,8))
sns.countplot(x='YearsAtCompany',hue = 'Attrition',data = medical_df)
plt.ylabel('Number of Employees')
plt.title('Number of Employees working years in the Company')
plt.show()
```



```
In [45]: from scipy.stats import levene
from scipy.stats import ttest_ind

print(levene(medical_df['YearsAtCompany'][medical_df['Attrition'] == 'Yes'],
            medical_df['YearsAtCompany'][medical_df['Attrition'] == 'No'], center = 'mean'))

print(ttest_ind(medical_df['YearsAtCompany'][medical_df['Attrition'] == 'Yes'],
            medical_df['YearsAtCompany'][medical_df['Attrition'] == 'No']))

# p < 0.05
# YearsAtCompany is significantly affect the attrition rate
# Employees in their early years of career tend to resign
# a lot more than employees who've worked in the company for more 7 years

LeveneResult(statistic=2.5049476586425445, pvalue=0.11370465713766563)
Ttest_indResult(statistic=-5.1963086670254235, pvalue=2.3188716103863033e-07)
```


Describe the Attrition based on Income. Compute the attrition rate based on Income

```
In [46]: mean_income = round(medical_df.groupby('Attrition',as_index=False)['MonthlyIncome'].mean(),2)
mean_income = mean_income.rename(columns={'MonthlyIncome': 'Mean Monthly Income'})
mean_income
```

```
Out[46]:
```

	Attrition	Mean Monthly Income
0	No	6832.74
1	Yes	4787.09

```
In [47]: sns.boxplot(x='Attrition', y='MonthlyIncome', data=medical_df)
plt.show()
```



```
In [48]: rate_att=medical_df.groupby(['MonthlyIncome', 'Attrition']).apply(lambda x:x['MonthlyIncome'].count()).reset_index(name='Counts')
rate_att['MonthlyIncome']=round(rate_att['MonthlyIncome'],-3)
rate_att=rate_att.groupby(['MonthlyIncome', 'Attrition']).apply(lambda x:x['MonthlyIncome'].count()).reset_index(name='Counts')
fig=px.line(rate_att,x='MonthlyIncome',y='Counts',color='Attrition',title='Monthly Income counts of People in Organization')
fig.show()
```


Monthly Income counts of People in Organization



```
In [49]: print(levene(medical_df['MonthlyIncome'][medical_df['Attrition'] == 'Yes'],
    medical_df['MonthlyIncome'][medical_df['Attrition'] == 'No'], center = 'mean'))

print(ttest_ind(medical_df['MonthlyIncome'][medical_df['Attrition'] == 'Yes'],
    medical_df['MonthlyIncome'][medical_df['Attrition'] == 'No'], equal_var = False))

# p < 0.05
# Monthly Income is significantly affect the attrition rate
# Higher salary is a usually motivation to keep working in the same company

LeveneResult(statistic=26.31302358992364, pvalue=3.29070575711061e-07)
Ttest_indResult(statistic=-7.482621586644742, pvalue=4.433588628286071e-13)
```

Describe the Attrition based on Environment. Compute the attrition rate based on Environment

```
In [50]: round(pd.crosstab(index=medical_df['Attrition'], columns=medical_df['EnvSatisfaction'], normalize='index'), 2)
```

```
Out[50]: EnvSatisfaction    1    2    3    4
Attrition
No      0.17  0.20  0.32  0.31
Yes     0.30  0.18  0.26  0.25
```

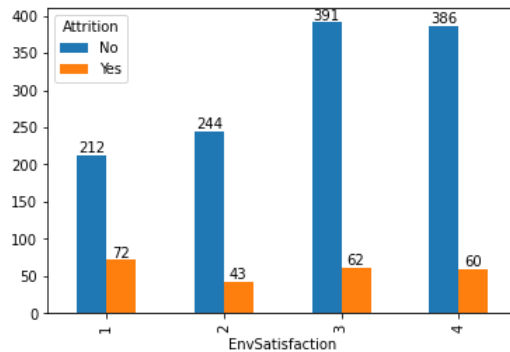
```
In [51]: envi_raw = pd.crosstab(index=medical_df['Attrition'], columns=medical_df['EnvSatisfaction'])
envi_raw
```

```
Out[51]: EnvSatisfaction    1    2    3    4
```

		Attrition			
	No	212	244	391	386
	Yes	72	43	62	60

```
In [52]: envi = pd.crosstab(index=medical_df['EnvSatisfaction'], columns=medical_df['Attrition'])
ax = envi.plot(kind = 'bar')

# annotate
ax.bar_label(ax.containers[0], label_type='edge')
ax.bar_label(ax.containers[1], label_type='edge')
plt.show()
```



```
In [53]: chi,p,dof,exp = stats.chi2_contingency(envi)
print('the chi square test is {}'.format(p))
# 5.1234689062894205e-05 < 0.05
# the satisfaction level from level 1-4, the chances of people leaving the company is slightly decrease.
# There is a change the company to improve and do better to retain the employees

the chi square test is 5.1234689062894205e-05
```

Describe the Attrition based on Job Satisfaction. Compute the attrition rate based on Job Satisfaction

```
In [54]: round(pd.crosstab(index=medical_df['Attrition'], columns=medical_df['JobSatisfaction'], normalize='index'),2)
```

```
Out[54]: JobSatisfaction    1    2    3    4
```

		Attrition			
	No	0.18	0.19	0.30	0.33
	Yes	0.28	0.19	0.31	0.22

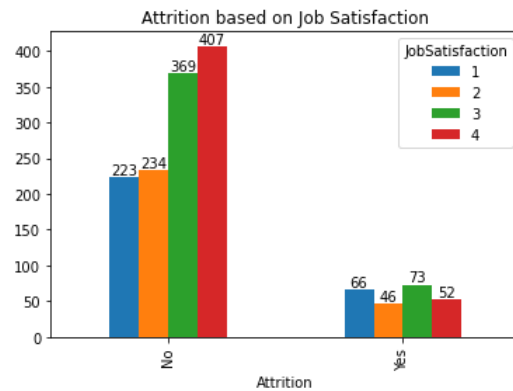
```
In [55]: satisfaction_raw = round(pd.crosstab(index=medical_df['Attrition'], columns=medical_df['JobSatisfaction']),2)
satisfaction_raw
```

Out[55]: **JobSatisfaction** 1 2 3 4

Attrition				
No	223	234	369	407
Yes	66	46	73	52

```
In [56]: ax = satisfaction_raw.plot(kind = 'bar')

# annotate
ax.bar_label(ax.containers[0], label_type='edge')
ax.bar_label(ax.containers[1], label_type='edge')
ax.bar_label(ax.containers[2], label_type='edge')
ax.bar_label(ax.containers[3], label_type='edge')
plt.title('Attrition based on Job Satisfaction')
plt.show()
```



```
In [57]: chi,p,dof,exp = stats.chi2_contingency(satisfaction_raw)
print('the chi square test is {}'.format(p))

# 0.0005563004510387556 < 0.05
# the job satisfaction level from Level 1-4, affect the attrition rate

# the chances of people leaving the company is slightly decrease.
# There is a change the company to improve and do better to retain the employees

the chi square test is 0.0005563004510387556
```


Describe the Attrition based on Age. Compute the attrition rate based on Age

```
In [58]: mean_age = round(medical_df.groupby('Attrition',as_index=False)['Age'].mean(),1)
mean_age = mean_age.rename(columns={'Age': 'Mean Age'})
mean_age
```


Out[58]:

	Attrition	Mean Age
0	No	37.6
1	Yes	33.6


```
In [59]: fig = plt.figure(figsize=(18,8))
sns.countplot(x='Age',hue = 'Attrition',data = medical_df)
plt.show()
```



```
In [60]: fig,ax = plt.subplots(figsize=(15,7))
medical_df.groupby(['Age','Attrition']).count()['Gender'].unstack().plot(ax=ax)
ax.set_title('Attrition based on Age in the Company')
# ax.grid()
plt.xticks(medical_df['Age'])
plt.ylabel('Number of Employees')
plt.show()
# the attrition is maximum between the age groups 28-32.
# The attrition rate keeps on falling with increasing age,
# as people look after stability in their jobs at these point of times
```



```
In [61]: sns.boxplot(x='Attrition', y='Age', data=medical_df)
plt.show()
```



```
In [62]: print(levene(medical_df['Age'][medical_df['Attrition'] == 'Yes'],
                    medical_df['Age'][medical_df['Attrition'] == 'No'], center='mean'))

print(ttest_ind(medical_df['Age'][medical_df['Attrition'] == 'Yes'],
                medical_df['Age'][medical_df['Attrition'] == 'No']))

# p < 0.05
# Age is significantly affect the attrition rate
# Older people is tend to keep working and stay in the same company

LeveneResult(statistic=1.1583170677572885, pvalue=0.2819916793250208)
Ttest_indResult(statistic=-6.1786638353072165, pvalue=8.356308021103649e-10)
```

Describe the Attrition based on BusinessTravel. Compute the attrition rate based on BusinessTravel

```
In [63]: round(pd.crosstab(index=medical_df['Attrition'],columns=medical_df['BusinessTravel'],normalize='index'),2)
```

```
Out[63]: BusinessTravel  Non-Travel  Travel_Frequently  Travel_Rarely
```

Attrition				
	No	Yes		
BusinessTravel	0.11	0.05	0.17	0.29
Non-Travel	0.72	0.66		

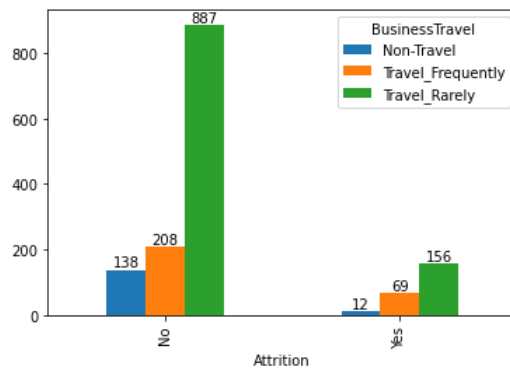
```
In [64]: travel_raw = round(pd.crosstab(index=medical_df['Attrition'],columns=medical_df['BusinessTravel']),2)
travel_raw
```

```
Out[64]: BusinessTravel  Non-Travel  Travel_Frequently  Travel_Rarely
```

Attrition				
	No	Yes		
BusinessTravel	138	12	208	69
Non-Travel	887	156		

```
In [65]: ax = travel_raw.plot(kind = 'bar')
```

```
# annotate
ax.bar_label(ax.containers[0], label_type='edge')
ax.bar_label(ax.containers[1], label_type='edge')
ax.bar_label(ax.containers[2], label_type='edge')
plt.show()
```



```
In [66]: chi,p,dof,exp = stats.chi2_contingency(travel_raw)
print('the chi square test is {}'.format(p))
# p < 0.05
# BusinessTravel is significantly affect the attrition rate
```

```
the chi square test is 5.608614476449931e-06
```

Describe the Attrition based on Overtime. Compute the attrition rate based on Overtime

```
In [67]: round(pd.crosstab(index=medical_df['Attrition'],columns=medical_df['OverTime'],normalize='index'),2)
```

Out[67]:

Attrition		
	No	Yes
No	0.77	0.23
Yes	0.46	0.54


```
In [68]: over_raw = round(pd.crosstab(index=medical_df['Attrition'], columns=medical_df['OverTime']), 2)
over_raw
```

Out[68]:

Attrition		
	No	Yes
No	944	289
Yes	110	127

```
In [69]: ax = over_raw.plot(kind = 'bar')
```

```
# annotate
ax.bar_label(ax.containers[0], label_type='edge')
ax.bar_label(ax.containers[1], label_type='edge')
plt.show()
```



```
In [70]: chi, p, dof, exp = stats.chi2_contingency(over_raw)
print('the chi square test is {}'.format(p))
# p < 0.05
# OverTime is significantly affect the attrition rate

the chi square test is 8.15842372153832e-21
```

```
In [71]: # marry_raw = round(pd.crosstab(index=medical_df['Attrition'], columns=medical_df['MaritalStatus']), 2)
# marry_raw
```

```
In [72]: # chi, p, dof, exp = stats.chi2_contingency(marry_raw)
# p
```

```
In [73]: # rearranging the columns
medical_df.columns
```

```
Out[73]: Index(['Age', 'BusinessTravel', 'MonthlyIncome', 'JobSatisfaction', 'Bonus',
      'Department', 'DistanceFromHome', 'Education', 'EducationField',
      'EnvSatisfaction', 'Gender', 'JobRole', 'MaritalStatus',
      'PerformanceRating', 'TrainingTimesLastYear', 'YearsAtCompany',
      'YearsSinceLastPromotion', 'OverTime', 'Attrition'],
      dtype='object')
```

```
In [74]: # medical_df = medical_df[['Age', 'BusinessTravel', 'MonthlyIncome', 'JobSatisfaction', 'Bonus',
#      'DistanceFromHome', 'EnvSatisfaction', 'YearsAtCompany',
#      'YearsSinceLastPromotion', 'OverTime', 'Attrition']]

#Logistic - overall 0.86 No      0.86 Yes      0.78
#naive - overall 0.82 No      0.87 Yes      0.41

# medical_df = medical_df[['Age', 'BusinessTravel', 'MonthlyIncome', 'JobSatisfaction', 'Bonus',
#      'Department', 'DistanceFromHome',
#      'EnvSatisfaction', 'JobRole',
#      'PerformanceRating', 'TrainingTimesLastYear', 'YearsAtCompany',
#      'YearsSinceLastPromotion', 'OverTime', 'Attrition']]
#Logistic - overall 0.86 No      0.86 Yes      0.86
#naive - overall 0.7 No      0.91 Yes      0.29

# medical_df = medical_df[['Age', 'BusinessTravel', 'MonthlyIncome', 'JobSatisfaction', 'DistanceFromHome',
#      'EnvSatisfaction', 'YearsSinceLastPromotion', 'OverTime', 'Attrition']]

#Logistic - overall 0.85 No 0.85      Yes 0.83
#naive - overall 0.84 No 0.86      Yes 0.47

# medical_df = medical_df[['Age', 'MonthlyIncome',
#      'EnvSatisfaction', 'YearsSinceLastPromotion', 'OverTime', 'Attrition']]

#Logistic - overall 0.85 No 0.85      Yes 0.75
#naive - overall 0.85 No 0.86      Yes 0.57

# medical_df = medical_df[['JobSatisfaction', 'DistanceFromHome', 'MaritalStatus', 'OverTime', 'BusinessTravel',
#      'PerformanceRating', 'TrainingTimesLastYear', 'Attrition']]

#Logistic - overall 0.84 No 0.85      Yes 0.57
#naive - overall 0.80 No 0.85      Yes 0.23

# 'Age', 'BusinessTravel', 'MonthlyIncome', 'JobSatisfaction', 'Bonus',
#      'DistanceFromHome', 'EnvSatisfaction', 'Gender', 'MaritalStatus',
#      'PerformanceRating', 'TrainingTimesLastYear', 'YearsAtCompany',
#      'YearsSinceLastPromotion', 'OverTime', 'Attrition'

# medical_df = medical_df[['JobSatisfaction', 'Age', 'OverTime', 'YearsAtCompany',
#      'EnvSatisfaction', 'TrainingTimesLastYear', 'Attrition']]

#Logistic - overall 0.85 No 0.85      Yes 0.62 f1-score no 0.91 yes 0.19
#naive - overall 0.85 No 0.85      Yes 0.62 f1-score no 0.92 yes 0.23

# medical_df = medical_df[['BusinessTravel', 'JobSatisfaction',
#      'DistanceFromHome', 'EnvSatisfaction', 'Gender', 'MaritalStatus',
#      'TrainingTimesLastYear', 'YearsAtCompany',
#      'YearsSinceLastPromotion', 'OverTime', 'Attrition']]

#Logistic - overall 0.85 No 0.86      Yes 0.67 f1-score no 0.92 yes 0.30
#naive - overall 0.80 No 0.86      Yes 0.33 f1-score no 0.88 yes 0.27

medical_df = medical_df[['Age', 'BusinessTravel', 'JobSatisfaction',
      'DistanceFromHome', 'EnvSatisfaction', 'Gender', 'MaritalStatus',
      'TrainingTimesLastYear', 'YearsAtCompany',
```



```
'YearsSinceLastPromotion', 'OverTime', 'Attrition']]

#Logistic - overall 0.87    No 0.87      Yes 0.78 f1-score no 0.9 yes 0.38
#naive - overall 0.82      No 0.87      Yes 0.40 f1-score no 0.9 yes 0.32

# medical_df = medical_df[['MonthlyIncome', 'BusinessTravel', 'JobSatisfaction',
# 'DistanceFromHome', 'EnvSatisfaction', 'Gender', 'MaritalStatus',
# 'TrainingTimesLastYear', 'YearsAtCompany',
# 'YearsSinceLastPromotion', 'OverTime', 'Attrition']]

#Logistic - overall 0.86    No 0.87      Yes 0.75 f1-score no 0.92 yes 0.33
#naive - overall 0.81      No 0.87      Yes 0.39 f1-score no 0.89 yes 0.33

# medical_df = medical_df[['BusinessTravel', 'JobSatisfaction',
# 'DistanceFromHome', 'EnvSatisfaction', 'Gender', 'MaritalStatus',
# 'TrainingTimesLastYear', 'YearsAtCompany',
# 'YearsSinceLastPromotion', 'OverTime', 'Attrition']]

# medical_df = medical_df[['Age', 'BusinessTravel', 'JobSatisfaction',
# 'Department', 'DistanceFromHome', 'Education', 'EducationField',
# 'EnvSatisfaction', 'Gender', 'JobRole', 'MaritalStatus',
# 'TrainingTimesLastYear', 'YearsAtCompany',
# 'YearsSinceLastPromotion', 'OverTime', 'Attrition']]

medical_df
```

Out[74]:

	Age	BusinessTravel	JobSatisfaction	DistanceFromHome	EnvSatisfaction	Gender	MaritalStatus	TrainingTimesLastYear	YearsAtCompany	YearsSinceLastPromotion	OverTime	Attrition
EmployeeID												
1	41	Travel_Rarely	4	1	2	Female	Single	0	6	0	Yes	Yes
2	49	Travel_Frequently	2	8	3	Male	Married	3	10	1	No	No
4	37	Travel_Rarely	3	2	4	Male	Single	3	0	0	Yes	Yes
5	33	Travel_Frequently	3	3	4	Female	Married	3	8	3	Yes	No
7	27	Travel_Rarely	2	2	1	Male	Married	3	2	2	No	No
...
2061	36	Travel_Frequently	4	23	3	Male	Married	3	5	0	No	No
2062	39	Travel_Rarely	1	6	4	Male	Married	5	7	1	No	No
2064	27	Travel_Rarely	2	4	2	Male	Married	0	6	0	Yes	No
2065	49	Travel_Frequently	2	2	4	Male	Married	3	9	0	No	No
2068	34	Travel_Rarely	3	8	2	Male	Married	3	4	1	No	No

1470 rows × 12 columns

Assigning X and y variable

```
In [75]: # Assign X and y values for the partitioning
X = medical_df.iloc[:, :-1]
y = medical_df.iloc[:, -1]
```

In [76]: X

Out[76]:

	Age	BusinessTravel	JobSatisfaction	DistanceFromHome	EnvSatisfaction	Gender	MaritalStatus	TrainingTimesLastYear	YearsAtCompany	YearsSinceLastPromotion	OverTime
EmployeeID											
1	41	Travel_Rarely	4	1	2	Female	Single	0	6	0	Yes
2	49	Travel_Frequently	2	8	3	Male	Married	3	10	1	No
4	37	Travel_Rarely	3	2	4	Male	Single	3	0	0	Yes
5	33	Travel_Frequently	3	3	4	Female	Married	3	8	3	Yes
7	27	Travel_Rarely	2	2	1	Male	Married	3	2	2	No
...
2061	36	Travel_Frequently	4	23	3	Male	Married	3	5	0	No
2062	39	Travel_Rarely	1	6	4	Male	Married	5	7	1	No
2064	27	Travel_Rarely	2	4	2	Male	Married	0	6	0	Yes
2065	49	Travel_Frequently	2	2	4	Male	Married	3	9	0	No
2068	34	Travel_Rarely	3	8	2	Male	Married	3	4	1	No

1470 rows × 11 columns

In [77]:

y

Out[77]:

```

EmployeeID
1      Yes
2      No
4      Yes
5      No
7      No
...
2061    No
2062    No
2064    No
2065    No
2068    No
Name: Attrition, Length: 1470, dtype: object

```

Handling Categorical Data

In [78]:

```

# Only transform the input variable
X = pd.get_dummies(X, drop_first=True)
X.head(5)

```

Out[78]:

	Age	JobSatisfaction	DistanceFromHome	EnvSatisfaction	TrainingTimesLastYear	YearsAtCompany	YearsSinceLastPromotion	BusinessTravel_Travel_Frequently	BusinessTravel_Travel_Rarely	Gender_Male	MaritalStatus
EmployeeID											
1	41	4	1	2	0	6	0	0	1	0	
2	49	2	8	3	3	10	1	1	0	1	
4	37	3	2	4	3	0	0	0	1	1	
5	33	3	3	4	3	8	3	1	0	0	
7	27	2	2	1	3	2	2	0	1	1	

Partitioning

In [79]:

```
# Partitioning the dataset
from sklearn.model_selection import train_test_split

X_train, X_test, y_train, y_test = train_test_split(X,y,test_size=0.3,
                                                    random_state=0, stratify = y) # using test size 0.3
```

In [80]: X_train

Out[80]:

	Age	JobSatisfaction	DistanceFromHome	EnvSatisfaction	TrainingTimesLastYear	YearsAtCompany	YearsSinceLastPromotion	BusinessTravel_Travel_Frequently	BusinessTravel_Travel_Rarely	Gender_Male	MaritalStatus
EmployeeID											
1523	44	1	28	4	1	20	14	0	1	1	
1212	40	2	2	2	2	0	0	1	0	1	
110	34	2	1	1	2	5	1	0	1	1	
615	40	1	28	3	2	1	0	0	0	1	
1161	37	4	25	3	2	6	1	0	1	0	
...
1061	50	3	2	2	2	8	7	0	1	1	
527	25	4	4	2	3	5	1	0	1	0	
602	39	3	8	3	3	8	0	1	0	0	
807	40	2	2	3	2	8	3	1	0	0	
689	37	3	3	3	2	10	7	0	1	1	

1029 rows × 13 columns

In [81]: print(len(X_train),len(X_test))

1029 441

In [82]: X_train

Out[82]:

	Age	JobSatisfaction	DistanceFromHome	EnvSatisfaction	TrainingTimesLastYear	YearsAtCompany	YearsSinceLastPromotion	BusinessTravel_Travel_Frequently	BusinessTravel_Travel_Rarely	Gender_Male	MaritalStatus
EmployeeID											
1523	44	1	28	4	1	20	14	0	1	1	
1212	40	2	2	2	2	0	0	1	0	1	
110	34	2	1	1	2	5	1	0	1	1	
615	40	1	28	3	2	1	0	0	0	1	
1161	37	4	25	3	2	6	1	0	1	0	
...	
1061	50	3	2	2	2	8	7	0	1	1	
527	25	4	4	2	3	5	1	0	1	0	
602	39	3	8	3	3	8	0	1	0	0	
807	40	2	2	3	2	8	3	1	0	0	
689	37	3	3	3	2	10	7	0	1	1	

1029 rows × 13 columns

In [83]: y_train

Out[83]:

EmployeeID	
1523	No
1212	No
110	No
615	No
1161	No
..	
1061	No
527	No
602	No
807	No
689	No

Name: Attrition, Length: 1029, dtype: object

Scaling the Data or perform Normalization

In [84]:

```
#Scaling the data
from sklearn.preprocessing import StandardScaler # Zscore normalization

scaler = StandardScaler()
scaler.fit(X_train) # compute the scalar model
X_train = scaler.transform(X_train) # here it will transform the X train into standard data
```

In [85]: X_train

```
Out[85]: array([[ 0.75909679, -1.54871647,  2.3244173 , ...,  1.09485978,
        -0.68097495, -0.63245553],
       [ 0.32796082, -0.64118273, -0.88376782, ...,  1.09485978,
        -0.68097495, -0.63245553],
       [-0.31874314, -0.64118273, -1.00715956, ...,  1.09485978,
        -0.68097495, -0.63245553],
       ...,
       [ 0.22017682,  0.26635101, -0.14341741, ...,  1.09485978,
        -0.68097495, -0.63245553],
       [ 0.32796082, -0.64118273, -0.88376782, ...,  1.09485978,
        -0.68097495, -0.63245553],
       [ 0.00460884,  0.26635101, -0.76037609, ...,  1.09485978,
        -0.68097495, -0.63245553]])
```

```
In [86]: # x_train_copy = pd.DataFrame(X_train)
        # x_train_copy
```

Train the Model using Logistic Regression

```
In [87]: #Train Logistic Regression Algorithm

        #Import LogisticRegression() function
        from sklearn.linear_model import LogisticRegression

        #Create a Logistic Regression classifier
        # random state is 0 because to maintain the data generated value
        medical_model = LogisticRegression(random_state=0)

        #Train the Logistic Regression algorithm using train set
        medical_model.fit(X_train,y_train)
```

```
Out[87]: LogisticRegression
```

Train the Model using Naive Bayes Thereom

```
In [88]: # # #Train Naive Bayes Thereom Algorithm
        # from sklearn.naive_bayes import GaussianNB

        # medical_model = GaussianNB()

        # medical_model.fit(X_train,y_train)
```

Apply the model to test set

```
In [89]: # Apply the model with the test set
        #Feature Scaling the Test Set
        from sklearn.preprocessing import StandardScaler
        X_test = scaler.transform(X_test)
        X_test
```

```
In [90]: # Apply for Test Set
y_pred=medical_model.predict(X_test)
y_pred
```

```
In [91]: #Dataframe of Predicted Output and Actual Output for Test set
df_validate = pd.DataFrame({'Actual':y_test,'Predicted':y_pred})
```

```
df_validate['Predicted']=df_validate['Predicted']
df_validate.head(5)
```

Out[91]:

	Actual	Predicted
EmployeeID		
916	No	No
1408	No	No
1233	No	No
1640	No	No
1133	No	No

Performance Evaluation

In [92]: *# use crosstab to evaluate the categorical*
 pd.crosstab(y_pred,y_test)

Out[92]:

	Attrition	No	Yes
row_0			
No	365	53	
Yes	5	18	

Precision & Sensitivity/Recall Report

- Logistic Regression

```
In [93]: 1 # Precision & Sensitivity/Recall Report
          2 from sklearn.metrics import classification_report
          3 print(classification_report(y_true=y_test,y_pred=y_pred))
```

	precision	recall	f1-score	support
No	0.87	0.99	0.93	370
Yes	0.78	0.25	0.38	71
accuracy			0.87	441
macro avg	0.83	0.62	0.65	441
weighted avg	0.86	0.87	0.84	441

- using Age -

Precision & Sensitivity/Recall Report

```
In [143]: 1 # Precision & Sensitivity/Recall Report
          2 from sklearn.metrics import classification_report
          3 print(classification_report(y_true=y_test,y_pred=y_pred))
```

	precision	recall	f1-score	support
No	0.87	0.99	0.92	370
Yes	0.75	0.21	0.33	71
accuracy			0.86	441
macro avg	0.81	0.60	0.63	441
weighted avg	0.85	0.86	0.83	441

- using Monthly Income -

```
In [89]: 1 # Precision & Sensitivity/Recall Report
          2 from sklearn.metrics import classification_report
          3 print(classification_report(y_true=y_test,y_pred=y_pred))
```

	precision	recall	f1-score	support
No	0.86	0.98	0.92	370
Yes	0.67	0.20	0.30	71
accuracy			0.85	441
macro avg	0.77	0.59	0.61	441
weighted avg	0.83	0.85	0.82	441

- using Non Monthly Income and Age -
- Naive Bayes

```
In [89]: 1 # Precision & Sensitivity/Recall Report
          2 from sklearn.metrics import classification_report
          3 print(classification_report(y_true=y_test,y_pred=y_pred))
```

	precision	recall	f1-score	support
No	0.87	0.92	0.90	370
Yes	0.40	0.27	0.32	71
accuracy			0.82	441
macro avg	0.64	0.60	0.61	441
weighted avg	0.79	0.82	0.80	441

- using Age -


```
In [93]: # Precision & Sensitivity/Recall Report
from sklearn.metrics import classification_report
print(classification_report(y_true=y_test,y_pred=y_pred))
```

	precision	recall	f1-score	support
No	0.87	0.99	0.93	370
Yes	0.78	0.25	0.38	71
accuracy			0.87	441
macro avg	0.83	0.62	0.65	441
weighted avg	0.86	0.87	0.84	441

```
In [ ]:
```